Aftertalk

BUILD BRIEF · APRIL 28, 2026 · DAYS 4–5
github.com/theaayushstha1/aftertalk

DAYS 4–5 OF 7 · 5 OF 7 STRETCHES SHIPPED · 2 DAYS TO DEADLINE

A neural voice, who said what, and a *quiet, editorial interface*.

Days 4 and 5 took the working voice loop and made it feel like a product. The phone now answers in a real human voice — not the system synthesizer — and labels who said what during the meeting. The interface was rebuilt around an editorial typography system, and a new Settings tab proves the privacy claim by reading live device state.

What shipped

COMMIT HASHES LINK TO GITHUB

Day 4

VOICE + SPEAKERS

Neural TTS and speaker diarization, both on-device. Swapped the placeholder system voice for **FluidAudio Kokoro 82M** running on the Neural Engine, with a sentence-boundary streamer that starts speaking before the LLM finishes generating. Wired **Pyannote 3.1 + WeSpeaker v2** in parallel with a Parakeet polish pass; the summary now attributes decisions and action items to specific speakers. iPhone Air stays under the jetsam ceiling via lazy-warm.

`6bee46d` `079e7d0` `8b594af` `3e63bee` `db56dbc` `c718264` `fd81f51`

Day 5

QUIET STUDIO + SETTINGS

Editorial UI refactor + cross-meeting memory wired to the surface + a privacy audit you can read. Every screen rebuilt against shared primitives (eyebrow, title, body, privacy badge, breathing orb, immersive waveform) with a warm light/dark palette. Global Ask routes through the hierarchical retriever; citations carry meeting title across meetings, timestamp within. Settings reads live `@Query` counts, probes model bundles via `FileManager`, watches `NWPathMonitor`, and exposes a Verify button — nothing on that screen is cached or hand-waved. First on-device install surfaced five UI bugs; all five patched the same evening.

`d1f7be5` `3e096af` `45a696f` `73f59db` `ebf82d1` `e87387b` `2448ef3`

The Quiet Studio interface

14 SCREENS, ONE EDITORIAL SYSTEM

ONBOARDING · CAPTURE · PROCESSING · DETAIL · ASK · SETTINGS

Two choices that mattered

DECISION · DAY 4

Sentence-boundary streaming over chunked TTS

The LLM emits tokens in <33ms, but Kokoro wants whole sentences for natural prosody. Built a **SentenceBoundaryDetector** that consumes the snapshot stream and hands completed sentences to a **TTSWorker** actor. Sentence N+1 inference prefetches while sentence N plays. Result: first spoken word in **104 ms** while the rest of the answer is still being generated.

DECISION · DAY 5

Settings reads only real device state

Privacy is the entire pitch, so the audit screen had to be auditable. Every row pulls live state per render — `@Query` for meeting and chunk counts, `FileManager` probes via `ModelLocator` for bundle integrity, `PrivacyMonitor.state` for the network row. The **Verify** button runs an actual probe sequence, not a hardcoded green check. Recorded in **ADR-016**.

The numbers (held from the on-device run)

104 ms	1.44 s	5/7	82 m	0
FIRST SPOKEN WORD	FULL Q&A TURN	STRETCH GOALS SHIPPED	KOKORO PARAMS ON ANE	NETWORK REQUESTS